



training loss



next-token prediction accuracy



perplexity

